



QuanTool Benchmark Analysis: INT8 Dynamic Quantization


(a) Accuracy: FP32 vs INT8 Dynamic



(b) Latency Speedup with INT8 Dynamic



(c) Model Size Reduction with INT8 Dynamic



(d) Accuracy-Latency Trade-off (bubble size = compression ratio)

